



Ruby Fundamentals Day Three

Dan Klopp,
Coban Tun

Table of Contents

- REST API's
 - Introduction and Design
 - Ruby Sample
- Creating the REST API
- Advanced
 - Scope
 - Procs, Blocks and Lambdas [not covered]

REST

- REpresentational Stateful Transfer
- Popular examples
 - Openstack and AWS
 - Sensu and Nagios
- A very common stateless architecture

Our Focus

- We are concerned with implementing a simple API
- We are NOT concerned with proper design
- We are concerned with terminology
- We are NOT concerned with security

HTTP Verbs

- PUT
 - Place the entity at the specified resource
- DELETE
 - Delete the specified resource
- GET
 - Get the entity at the specified resource
- POST
 - Instruct resource to act on enclosed entity

HTTP PUT

- PUT
 - Place the entity at the specified URI
 - `curl -X PUT -d arg=val -d arg2=val2 localhost:8080`
- If entity already exists at location, it updates the entity
- Example
 - `PUT 95134 localhost/zipcode`

HTTP DELETE

- Delete the URI
- Note: successful return code does not mean resource was deleted, only that the server intends to delete it
- Example:
 - DELETE localhost/zipcode/95134

HTTP GET

- Get the entity at the specified resource
- Should be idempotent
- Should only retrieve
 - This is a “guideline”
 - <https://www.youtube.com/watch?v=jl0hMfqNQ-g>
- Example
 - GET localhost/zipcode/95134

<<https://www.ietf.org/rfc/rfc2616.txt>>

HTTP POST

- Submit the enclosed entity to the URI handler
- This is fundamentally different from PUT
 - PUT places entity at URI
 - POST asks URI to act on entity
- Message is encoded in HTTP body, not URI
- Example
 - POST `getallbusinessesat localhost/zipcode/95134`

POST in Practice

- Theoretically POST should only be used for changes

POST in Practice

- Theoretically POST should only be used for changes
- In practice also used as a GET

POST in Practice

- Theoretically POST should only be used for changes
- In practice also used as a GET
 - GET is limited to 4000 characters or less (implementation dependent)
 - POST has no such limit

We Will Be Using GET

- It makes it simple to implement

We Will Be Using GET

- It makes it simple to implement
- Strictly speaking we are violating the standard

We Will Be Using GET

- It makes it simple to implement
- Strictly speaking we are violating the standard
- Properly implementing HTTP requests and responses is beyond the scope of a fundamentals course

AWS HTTP GET Example

- Querying AWS for metadata

```
# curl -X GET http://169.254.169.254/latest/meta-data/  
ami-id  
ami-launch-index  
ami-manifest-path  
block-device-mapping/  
hostname  
instance-action
```

- Polling the hostname

```
# curl -X GET http://169.254.169.254/latest/meta-data/hostname  
ip-10-248-111-220.us-west-2.compute.internal
```


Example Server

- Live Demonstration
- Source code at:

https://github.com/dsklopp/rest_example_server.git

Building A REST API

- First, let's design the API
- GET /
 - returns all keys
- GET /key
 - returns value of key "key"
- GET /key/value
 - sets value of key "key" to value

Building A REST API

- Where to begin?

Building A REST API

- Where to begin?
 - We'll create a stub and incrementally add functionality

The Sample API

- First step

```
class SocketListen
  def initialize port=8181
    @data = { 'apache' => 'apache',
              'bsd' => 'mit', 'chef' => 'apache' }
    puts @data
  end
end
SocketListen.new()
```

The Sample API

- Output

```
: tyrion @ Casterly Rock ; ruby rest_1.rb  
{"apache"=>"apache", "bsd"=>"mit", "chef"=>"apache"}  
: tyrion @ Casterly Rock ;
```

The Sample API

- Output

```
: tyrion @ Casterly Rock ; ruby rest_1.rb  
{"apache"=>"apache", "bsd"=>"mit", "chef"=>"apache"}  
: tyrion @ Casterly Rock ;
```

- Let's add a socket

The Sample API

- We'll use the TCPSocket Class of the Socket Module

```
require 'socket'

class SocketListen
  def initialize port=8181
    @data = { 'apache' => 'apache',
              'bsd' => 'mit', 'chef' => 'apache' }
    @port = port || 8181
    @server = TCPServer.new @port
  end
end
SocketListen.new()
```


The Sample API

- Create an infinite loop to listen for connections after creating the TCPServer

```
@server = TCPServer.new @port
loop {
  client = @server.accept
  request = client.gets
  puts request
  client.close
}
```

The Sample API

- Running that code and calling it via:

```
curl -X GET localhost:8181
```

- Will show the output of:

```
GET / HTTP/1.1  
GET / HTTP/1.1  
GET / HTTP/1.1  
GET / HTTP/1.1
```

- See “rest_3.rb” for sample
- Live demonstration

The Sample API

- It would be nice if the server could talk with the client...

```
@server = TCPServer.new @port
loop {
  client = @server.accept
  request = client.gets
  puts request
  client.puts request
  client.close
}
```

- This will echo the client's request back to the client
- Live demonstration, see "ruby_3b.rb"

The Sample API

- When handling a request, we'll send the request to a dedicated method. We'll call the method “handle”.
- The call is:

```
@server = TCPServer.new @port
loop {
  client = @server.accept
  request = client.gets
  puts request
  client.puts handle request
  client.close
}
```

The Sample API

- For those more comfortable with parenthesis around method calls, this is equivalent code:

```
@server = TCPServer.new @port
loop {
  client = @server.accept
  request = client.gets()
  puts(request)
  client.puts(handle(request))
  client.close()
}
```

The Sample API

- The handle method only returns “GET” for the GET HTTP verb.

```
def handle arg
  verb = arg.split()[0]
  arg=arg.split()[1..-1].join(" ")
  if verb == 'GET'
    return "GET  "
  else
    return "NOT IMPLEMENTED"
  end
end
```

- Live demonstration, see “rest_4.rb”

The Sample API

- So far we have
 - A socket that listens for requests
 - A program that returns the request over the socket
 - Trivial request validation
- Next, act on GET URI requests

Building A REST API

- Refresher, remember we are responding to:
 - GET /
 - returns all keys
 - GET /key
 - returns value of key "key"
 - GET /key/value
 - sets value of key "key" to value

The Sample API

- Modify

```
def handle arg
  verb = arg.split()[0]
  arg=arg.split()[1..-1].join(" ")
  if verb == 'GET'
    return "GET    "
  else
    return "NOT IMPLEMENTED"
  end
end
```

- to...

The Sample API

- We are deferring action on the URI to another method

```
def handle arg
  verb = arg.split()[0]
  arg=arg.split()[1..-1].join(" ")
  if verb == 'GET'
    return handle_get arg
  else
    return "NOT IMPLEMENTED"
  end
end
```

The Sample API

- How do we parse the URI?

The Sample API

- How do we parse the URI?
- We'll convert the relevant portion into an array

The Sample API

- How do we parse the URI?
- We'll convert the relevant portion into an array

```
>> arg="GET /bsd HTTP/1.1"  
=> "GET /bsd HTTP/1.1"
```

The Sample API

- How do we parse the URI?
- We'll convert the relevant portion into an array

```
>> arg="GET /bsd HTTP/1.1"  
=> "GET /bsd HTTP/1.1"  
>> arg.split()  
=> ["GET", "/bsd", "HTTP/1.1"]
```

The Sample API

- How do we parse the URI?
- We'll convert the relevant portion into an array

```
>> arg="GET /bsd HTTP/1.1"
=> "GET /bsd HTTP/1.1"
>> arg.split()
=> ["GET", "/bsd", "HTTP/1.1"]
>> # We've already split the string earlier in handle()
?> # Handle did this and then passed it on to handle_get
?> arg=arg.split()[1..-1].join(" ")
=> "/bsd HTTP/1.1"
```

The Sample API

```
?> arg=arg.split()[1..-1].join(" ")  
=> "/bsd HTTP/1.1"
```


The Sample API

```
?> arg=arg.split()[1..-1].join(" ")  
=> "/bsd HTTP/1.1"  
>> # Isolate the URI  
?> arg.split()[0]  
=> "/bsd"
```

The Sample API

```
?> arg=arg.split()[1..-1].join(" ")  
=> "/bsd HTTP/1.1"  
>> # Isolate the URI  
?> arg.split()[0]  
=> "/bsd"  
>> # Break up each argument  
?> (arg.split[0]).split('/')  
=> ["", "bsd"]
```

The Sample API

```
?> arg=arg.split()[1..-1].join(" ")  
=> "/bsd HTTP/1.1"  
>> # Isolate the URI  
?> arg.split()[0]  
=> "/bsd"  
>> # Break up each argument  
?> (arg.split[0]).split('/')  
=> ["", "bsd"]  
>> # Remove the empty string before the first /  
?> (arg.split[0]).split('/')[1..-1]  
=> ["bsd"]
```

The Sample API

- First handle “GET /”

```
def handle_get arg
  path = (arg.split[0]).split('/')[1..-1]
  if path.nil?
    keys=[]
    @data.each do |key,value|
      keys << key
    end
    puts keys
    return keys
  end
end
```

The Sample API

- Now handle “GET /\$key”

```
else
    puts path[0] + " => " + @data[path[0]]
    return @data[path[0]] + "      "
```

The Sample API

- Now handle “GET /\$key/\$value”

```
elif path.size == 2
  @data[path[0].chomp] = path[1]
  puts path[1]
  return path[1]
```

The Sample API

- Now handle “GET /\$key”

```
def handle_get arg
  path = (arg.split[0]).split('/')[1..-1]
  if path.nil?
    keys=[]
    @data.each do |key,value|
      keys << key
    end
    puts keys
    return keys
  end
end
```

Scope

Code References

- The sample REST API for reporting licenses

https://github.com/dsklopp/rest_example_server

- A functional HAProxy Management tool

https://github.com/dsklopp/haproxy_rest_rubyclass

- All slide related code

https://github.com/dsklopp/ruby_class_2015_samples

In Class Exercise / Homework

- As demonstrated in class, add an API to your haproxy management tool.
 - The calls it must honor are:
 - GET /backend - show the active backend
 - GET /backends - show all available backends
 - GET /backend/\$backend - set the backend to \$backend and restart haproxy

In Class Exercise / Homework Comments

- This will not be easy
- Since the class is small we'll help the class out one on one as needed
- AWS instances will be available until Monday (Feb 2)

Questions?

